

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import r2_score, accuracy_score, mean_squared_error, mean_absolute_error
```

```
In [3]: df=pd.read_csv("HousingData.csv")
```

```
In [4]: df.head()
```

```
Out[4]:
```

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	B	LSTAT
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1	296	15.3	396.90	15.3
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2	242	17.8	396.90	17.8
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2	242	17.8	392.83	17.8
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3	222	18.7	394.63	18.7
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3	222	18.7	396.90	18.7

```
In [6]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 506 entries, 0 to 505
Data columns (total 14 columns):
#   Column      Non-Null Count  Dtype
---  -
0   CRIM        486 non-null    float64
1   ZN          486 non-null    float64
2   INDUS       486 non-null    float64
3   CHAS        486 non-null    float64
4   NOX         506 non-null    float64
5   RM          506 non-null    float64
6   AGE         486 non-null    float64
7   DIS         506 non-null    float64
8   RAD         506 non-null    int64
9   TAX         506 non-null    int64
10  PTRATIO     506 non-null    float64
11  B           506 non-null    float64
12  LSTAT       486 non-null    float64
13  MEDV        506 non-null    float64
dtypes: float64(12), int64(2)
memory usage: 55.5 KB
```

```
In [7]: df.shape
```

```
Out[7]: (506, 14)
```

```
In [8]: df.isnull().sum()
```

```
Out[8]: CRIM      20  
        ZN        20  
        INDUS    20  
        CHAS     20  
        NOX       0  
        RM        0  
        AGE      20  
        DIS       0  
        RAD       0  
        TAX       0  
        PTRATIO   0  
        B         0  
        LSTAT     20  
        MEDV      0  
        dtype: int64
```

```
In [9]: df=df.dropna()
```

```
In [10]: df.isnull().sum()
```

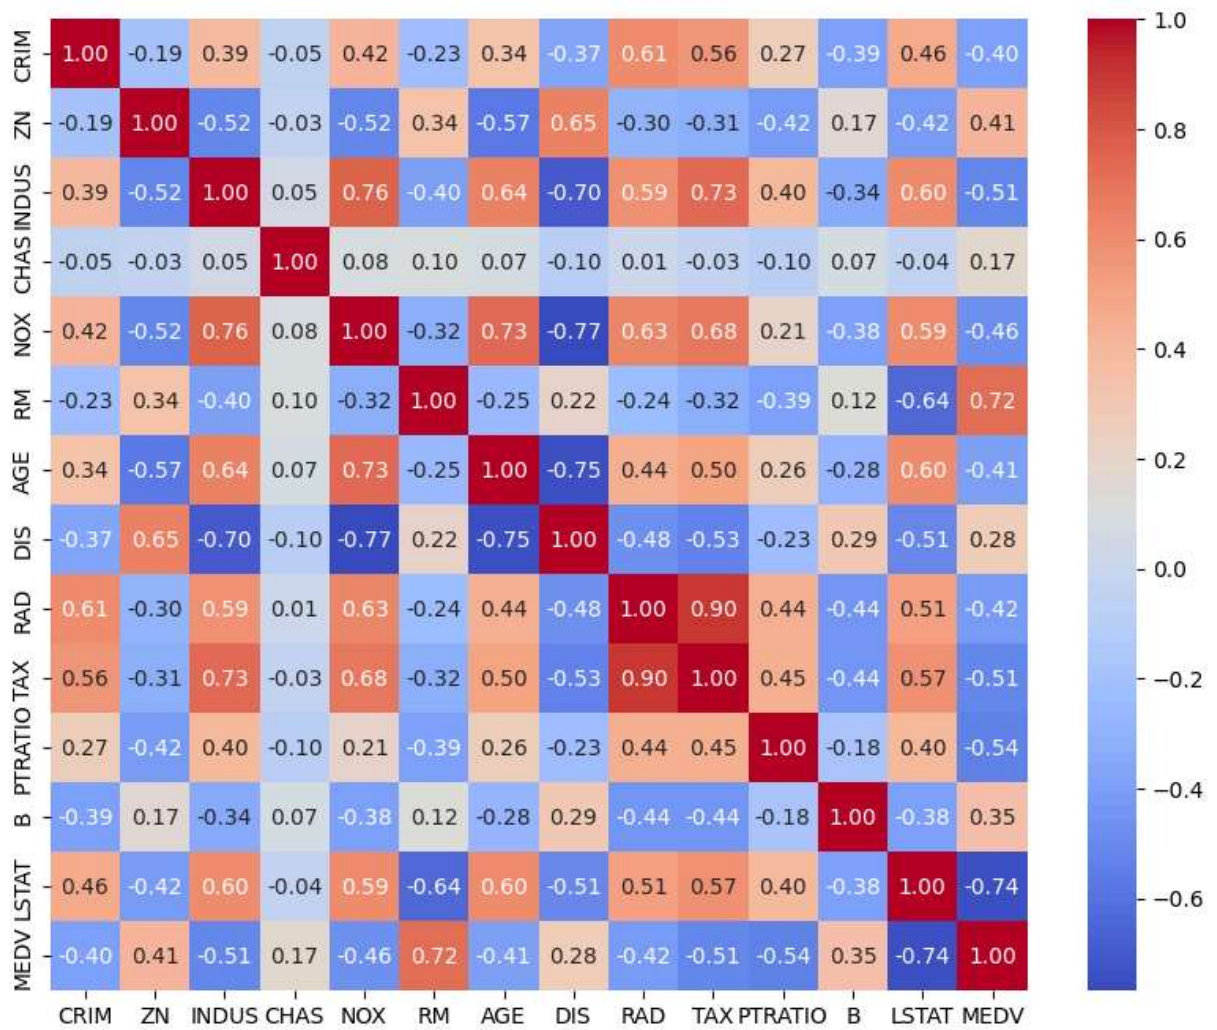
```
Out[10]: CRIM      0  
        ZN        0  
        INDUS     0  
        CHAS      0  
        NOX       0  
        RM        0  
        AGE       0  
        DIS       0  
        RAD       0  
        TAX       0  
        PTRATIO   0  
        B         0  
        LSTAT     0  
        MEDV      0  
        dtype: int64
```

```
In [12]: df.shape
```



```
Out[12]: (394, 14)
```

```
In [13]: corr=df.corr()
```

```
In [15]: plt.figure(figsize=(10, 8))  
        sns.heatmap(corr, annot=True, cmap="coolwarm", fmt=".2f")  
        plt.show()
```



```
In [17]: df.plot(x='LSTAT',y='MEDV',style='o')
plt.xlabel('lstat')
plt.ylabel('medv')
plt.show()
df.plot(x='RM',y='MEDV',style='o')
plt.xlabel('RM')
plt.ylabel('medv')
plt.show()
```



```
In [24]: df1=df.loc[:,['RM','MEDV']]
df1.head()
```

Out[24]:

	RM	MEDV
0	6.575	24.0
1	6.421	21.6
2	7.185	34.7
3	6.998	33.4
5	6.430	28.7

```
In [33]: x=pd.DataFrame(df1['RM'])
y=pd.DataFrame(df1['MEDV'])
```

```
In [34]: x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3,random_state=2)
```

```
In [35]: lr=LinearRegression()
lr.fit(x_train,y_train)
```

Out[35]:

▼ LinearRegression

LinearRegression()

```
In [43]: y_pred=lr.predict(x_test)
```

```
In [44]: lr.intercept_
```

Out[44]: array([-36.46388815])

```
In [45]: lr.coef_
```

Out[45]: array([[9.3646384]])

```
In [46]: y_pred.shape
```

Out[46]: (119, 1)

```
In [47]: y_test.shape
```


Out[47]: (119, 1)

```
In [48]: print("Mean Absolute Error :",mean_absolute_error(y_test,y_pred))
print("Meas Squared Error :", mean_squared_error(y_test,y_pred))
print("Root Mean Squared Error :", mean_squared_error(y_test,y_pred))
```

Mean Absolute Error : 4.21025512996237  
Meas Squared Error : 36.841866268951804  
Root Mean Squared Error : 36.841866268951804

```
In [52]: plt.scatter(y_test,y_pred)
plt.xlabel('y_test')
plt.ylabel('y_pred')
```

```
plt.show()
```



```
In [56]: from sklearn.preprocessing import StandardScaler
```

```
In [63]: sc=StandardScaler()
x_sc=sc.fit_transform(x)
x_train, x_test, y_train, y_test = train_test_split(x_sc, y, test_size=0.2, random_
```

```
In [64]: lr_scaled = LinearRegression()
lr_scaled.fit(x_train, y_train)
y_pred_scaled = lr_scaled.predict(x_test)

print("Scaled R²:", r2_score(y_test, y_pred_scaled))
```

Scaled R²: 0.4786797724382229

```
In [67]: from sklearn.preprocessing import PolynomialFeatures
pf=PolynomialFeatures()
x_poly=pf.fit_transform(x_sc)
plm=LinearRegression()
x_train, x_test, y_train, y_test = train_test_split(x_poly, y, test_size=0.2, random_
plm.fit(x_train, y_train)
y_pred_scaled = plm.predict(x_test)
```

```
In [68]: print("Scaled R²:", r2_score(y_test, y_pred_scaled))
```

Scaled R²: 0.5134161272669115

```
In [72]: input=[[8.95]]
         i=sc.transform(input)
         predict=lr_scaled.predict(i)
         predict
```

```
C:\Users\nitin\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\base.py:465: UserWarning: X does not have valid feature names, but StandardScaler was fitted with feature names
  warnings.warn(
```

```
Out[72]: array([[47.17608766]])
```

```
In [ ]:
```